

CISC/CMPE452/COGS400/CISC874

Unsupervised Learning I

Ch. 5 - Textbook

Farhana Zulkernine

Unsupervised Learning


- Small kids are able to recognize patterns => a significant amount of learning is accomplished by biological processes as “unsupervised” without a teacher.
- “Unsupervised” learning proceeds to discover special features and pattern from available data without using external help.

Unsupervised Learning Applications

- Tasks for which this is used are:
 - Clustering
 - Vector quantization
 - Approximation of data distribution
 - Feature extraction and
 - Dimensionality reduction
- Most unsupervised techniques bear resemblance to existing statistical methods.


Clustering

- Given a number of data points, determine a set of representative **centroids**, (or also called **prototypes**, **cluster centers**, or **reference vectors**)
- Find distances of a given pattern from the centroids.
- Typically it belongs to the cluster whose centroid is the closest.




centroids indicated in red


Example: Clustering



Reasonable number of clusters



Small Number of Clusters



Too Many Clusters

Fig. Three different ways of clustering the same set of sample points.

Clustering (cont...)


- Clusters are *evaluated* by measuring the average squared distance between each input pattern and the centroid of the cluster in which it is placed.

$$E_{\text{cluster}} = \frac{1}{\text{number of patterns}} \sum_{\text{patterns}} \|(\text{pattern} - \text{centroid})\|^2$$

$$E_{\text{total}} = \sum_{\text{clusters}} E_{\text{cluster}}$$

Clustering (cont...)

- In NN used for clustering, weights from the input layer to each output node constitute a **weight vector w** of that node, which represents the **centroid** of one cluster of input patterns.



Vector Quantization

- This is a task that applies unsupervised learning to **divide an input space into several connected regions** called *Voronoi Regions*, representing a quantization of the space.
- Each **region is represented** using a single vector called a *Codebook Vector (CV)* which are also called *Voronoi centres*.





Figure: **Voronoi Diagram** with 8 voronoi regions and codebook vectors of $\{cv_1, \dots, cv_8\}$

Vector Quantization


- Every point in the input space belongs to a region and is mapped to the corresponding CV.
 - Many input data vectors may be mapped to the same CV.
 - Two training approaches
 - Region may be given for each data point in a supervised learning and the CV has to be determined.
 - In unsupervised learning, a normal clustering approach can be applied.
- Therefore, the set of CVs is a **compressed form of information represented by all input data.**

Approximation of Data Distribution

- The data distribution in b is a concise description of the larger amount of data in a drawn from a *probability distribution*.
 - Data distribution in b is considered to be an approximation of that in a .
 - Points in b **may not be** a subset of points in a ← can be done using **unsupervised learning**.
 - Clustering, if used, will extract only one point for each cluster.



a




b

Feature Extraction and Dimensionality Reduction

- Patterns in different clusters should ideally be distinguished by some *feature*. Ex. All red balls in one cluster and green in the other – ‘color’ is the identifying feature.
- The goal of **feature extraction** is to find the most important features, i.e., those with the highest variation in a given population.
- Important **side-effect** is **reduction of input dimensionality** and thereby, improve in processing time and cost.

Winner-Take-All Networks

- Most unsupervised neural networks rely on **competitive learning** algorithms to compute and compare distances, determine the "winner" node with the highest level of activation, and use that to adapt weights.
- Patterns in the same cluster are as alike as possible.



Hamming Networks




Figure: A network to calculate Hamming distance (H) between stored vectors and input vectors

- **Weights** on links from an input layer to an output layer **represent components of stored input patterns**.
- Hamming networks **compute the "hamming distance"**, the number of differing bits, of input and stored vectors.
- So, with P output nodes a NN can store P vectors each associated with a weight vector.

Hamming Networks (cont...)

- Let $i_{p,j} \in \{+1, -1\}$ be the j th element of the p th stored vector.
- Components of the weight matrix W and threshold θ vector are given by $w_{p,j} = i_{p,j} / 2$, $j=1, \dots, n$ and $p=1, \dots, P$ and $\theta = -(n/2)$, respectively.
 - Uses constant bias.
- Input layer has n nodes (dimension of input vector) and output layer has P nodes (total number of stored vectors)
- p th output node generates the *negative of Hamming distance* between p th stored pattern and the input pattern.

Hamming Networks (cont...)

- When a new vector i is presented to this network, its upper level nodes generate the output as given below where $i_p \cdot i$ represents the dot product $\sum_k i_{p,k} i_k$
- *One shot training* \rightarrow Assign weights

$$W = \frac{1}{2} \begin{pmatrix} i_1^T \\ \vdots \\ i_P^T \end{pmatrix} \quad \Theta = \begin{pmatrix} -\frac{n}{2} \\ \vdots \\ -\frac{n}{2} \end{pmatrix} \quad o = Wi + \Theta = \frac{1}{2} \begin{pmatrix} i_1 \cdot i - n \\ \vdots \\ i_P \cdot i - n \end{pmatrix}$$

Example 5.1

- Given $i_1=(1, -1, -1, 1, 1)$, $i_2 = (-1, 1, -1, 1, -1)$, $i_3 = (1, -1, 1, -1, 1)$ and test data $x=(1, 1, 1, -1, -1)$, design a Hamming network.
- Weight matrix represents the 3 stored vectors i_1 , i_2 and i_3 .

$$W = \frac{1}{2} \begin{pmatrix} i_1^T \\ i_2^T \\ i_3^T \end{pmatrix} = \frac{1}{2} \begin{pmatrix} 1 & -1 & -1 & 1 & 1 \\ -1 & 1 & -1 & 1 & -1 \\ 1 & -1 & 1 & -1 & 1 \end{pmatrix} \quad \Theta = \frac{1}{2} \begin{pmatrix} -5 \\ -5 \\ -5 \end{pmatrix}$$

Example 5.1

$$1 - 1 - 1 - 1 - 1 = -3$$

- Given test data $x = (1, 1, 1, -1, -1)$, output can be calculated as: (See book for details)

$$Wx = \frac{1}{2} \begin{pmatrix} 1 & -1 & -1 & 1 & 1 \\ -1 & 1 & -1 & 1 & -1 \\ 1 & -1 & 1 & -1 & 1 \end{pmatrix} \begin{bmatrix} 1 \\ 1 \\ 1 \\ -1 \\ -1 \end{bmatrix} = \frac{1}{2} \begin{bmatrix} -3 \\ -1 \\ 1 \end{bmatrix}$$

$$o = Wx + \Theta = \frac{1}{2} \begin{bmatrix} -3 - 5 \\ -1 - 5 \\ 1 - 5 \end{bmatrix} = \begin{pmatrix} -4 \\ -3 \\ -2 \end{pmatrix}$$

x and i_1 differs by 4 bits, i_2 by 3 bits, i_3 by 2 bits


Hamming Networks (cont...)

- Interestingly if all bits match ($w_i = n/2$) then output Hamming distance = 0, otherwise < 0 .
- So we can determine *which stored pattern is nearest to a new input pattern by taking the maximum of the outputs*.
- This can be accomplished by attaching a *Maxnet* on top of the second layer of the Hamming network.

Maxnet

- A Maxnet is a **recurrent competitive one-layer** network used to determine which node has the highest initial activation.
- $\theta = 1$ and $\varepsilon \leq -1/(\# \text{ of nodes})$ ensures that the **node that has the initial highest value prevail** as "winner", while the others subside to zero.
- The node function is $f(\text{net}) = \max(0, \text{net})$ where $\text{net} = \sum_{i=1}^n w_i x_i$

Similar to inhibitory lateral connections in the IAM model.



Maxnet (cont...)

- All nodes update their outputs simultaneously. Each node receives *inhibitory* inputs from all other nodes, via "lateral" (intra-layer) connections.
- The maxnet allows for greater parallelism in execution, since every computation is local to each node rather than centralized.

Example 5.2

- Let initial activation values = $(0.5, 0.9, 1, 0.9, 0.9)$ and $\epsilon = -1/5$ and $\theta = +1$. Computing outputs o_j after 1st iteration

$$o_1 = \max(0, 0.5 - 1/5(0.9 + 1 + 0.9 + 0.9)) = \max(0, -0.24) = 0$$

$$o_2 = \max(0, 0.9 - 1/5(0.5 + 1 + 0.9 + 0.9)) = \max(0, 0.24) = 0.24$$

$$o_3 = \max(0, 1 - 1/5(0.5 + 0.9 + 0.9 + 0.9)) = \max(0, 0.36) = 0.36$$

- In subsequent iterations,
 $(0, 0.24, 0.36, 0.24, 0.24) \rightarrow (0, 0.072, 0.216, 0.072, 0.072) \rightarrow (0, 0, 0.1728, 0, 0)$
- 3rd node becomes the winner although others had values very close to this node.

Simple Competitive Learning

- The Hamming net and Maxnet assist more complex unsupervised learning networks, helping to determine the *node whose weight vector is nearest to an input pattern*.
- The figure shows a generalized version where inputs are n-dimensional real value vectors, $\mathbb{R}^n \rightarrow [0,1]$.
- Also known as **Kohonen Learning**.

$$\text{Output} = \begin{cases} 1 & \text{if node is a winner} \\ 0 & \text{otherwise} \end{cases}$$

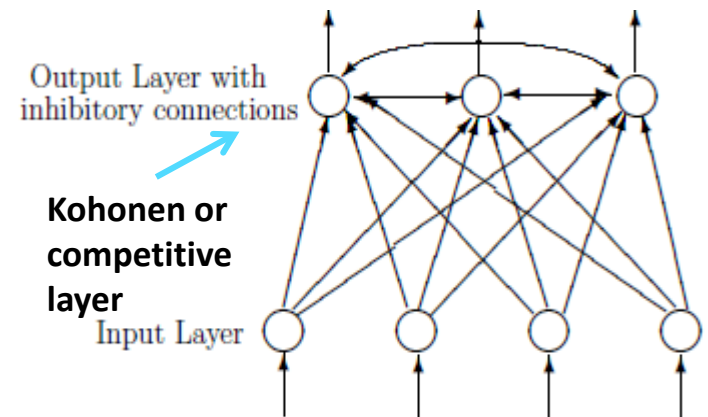



Figure: A simple competitive learning network

Design Objectives


- The objective is to **attach each output node to a stored pattern as represented by its weight vector.**
- Any **distance measure** at the output layer estimate how similar an input vector is to a weight vector.
 - Winner node at Kohonen layer is closest to the input vector.
- **Iterative training → Weight update rule:**
Adjust the weights of the winner output node such that w_i for Kohonen node i is as near as possible to all input samples for which the node is winner of the competition.

Clustering using Euclidean Distance




Example of competitive learning with three hidden nodes


Clustering (cont...)




Clustering (cont...)




Clustering (cont...)




Clustering (cont...)




Clustering (cont...)




Clustering (cont...)




Clustering (cont...)




Clustering (cont...)




Clustering (cont...)




Clustering (cont...)



Clustering (cont...)




Clustering (cont...)




- Continues for the epoch and for multiple iterations

Clustering (cont...)

- At the end the outputs of the network are at the center of the input data.
- Depending on their initial random values and the order of presentation of input data, the centroids of the clusters may be different.



Kohonen Learning (cont...)



- Simple competitive learning can be accomplished using a Maxnet.
- The j th output node is described by its weight vector from the input nodes,

$$w_j = (w_{j,1}, \dots, w_{j,n}) \quad l = \{1, \dots, n\} \text{ input dimension}$$

- A competition occurs to find the “winner” in the outer layer node j^* whose weight vector (or “prototype” or “node position”), is nearest to the input vector i_l
 $d(w_{j^*}, i_l) \leq d(w_j, i_l) \quad \text{where } j = \{1, \dots, m\} \text{ \# of output nodes}$
 $k = \{1, \dots, P\} \text{ \# of data points, } d \text{ is **Euclidean distance** =}$

$$d(w_j, i_k) = \sqrt{\sum_{l=1}^n (i_{l,k} - w_{j,l})^2}$$

How Kohonen Learning Works

- $d(w_j, i_l) = \sqrt{\sum_{l=1}^n (i_{l,k} - w_{j,l})^2}$ When expanded with $\|i_l\| = \|w_j\| = 1$
(unit vectors $\|p\| = \sqrt{p_1^2 + p_2^2 + \dots + p_n^2} = \sqrt{p \cdot p}$)
$$d^2(w, i) = i_k \cdot i_k + w_j \cdot w_j - 2 \sum i_{lk} w_{jl} = 2 - 2 \sum i_{lk} w_{jl} = 2 - 2 i \cdot w$$
- So, the distance will be minimum when $i \cdot w$ (dot product) is maximum ($\|i\| \|w\| \cos \theta$ is maximum at $\theta=0$). To generate output $i \cdot w$, use linear output neuron.
- *Hamming network* helps detect the distance as $i \cdot w$ in the first phase.
- Highest $i \cdot w$ indicates closest centroid. In the *winner-takes-all* phase, the *maxnet* recurrent network in the output layer finds the maximum distance $i \cdot w$ as the winning node.

Kohonen Learning (cont...)

- Weights w_j^* of the winning node j^* are adjusted keeping other weights unchanged so that w_j^* moves closer to the input i_l .

$$\Delta w_{j^*} = \eta (i_l - w_{j^*})$$

- It can be shown that in the limiting case, i.e., if and when **there is no significant change in the weight vector** when an input pattern is presented, $\Delta w_j < \varepsilon$, a predefined minimum threshold, w_j converges to

$$w_j = \frac{1}{\sum_{\ell} \delta_{j,\ell}} \sum_{\ell} i_{\ell} \delta_{j,\ell},$$
$$\delta_{j,\ell} = \begin{cases} 1 & \text{if the } j\text{th node is "winner" for input } i_{\ell}, \\ 0 & \text{otherwise,} \end{cases}$$

averaging input vectors for which w_j is the winner.

- The learning rate may vary such that $\eta(t+1) \leq \eta(t)$, resulting in faster convergence (gradually reduce the rate).

See example 5.3 in the book

Kohonen Learning (cont...)

Figure 5.4 Simple competitive learning algorithm

Initialize weights randomly;

repeat

- (Optional:) Adjust learning rate $\eta(t)$;
- Select an input pattern i_k ;
- Find node j^* whose weight vector w_{j^*} is closest to i_k ;
- Update each weight $w_{j^*,1}, \dots, w_{j^*,n}$ using the rule:

$$\Delta w_{j^*,l} = \eta(t)(i_{k,l} - w_{j^*,l}) \quad \text{for } \ell \in \{1, \dots, n\}$$

until network converges or computational bounds are exceeded

Example 5.3

Input vectors $T = \{i_1 = (1.1, 1.7, 1.8), i_2 = (0, 0, 0), i_3 = (0, 0.5, 1.5), i_4 = (1, 0, 0), i_5 = (0.5, 0.5, 0.5), i_6 = (1, 1, 1)\}$

The network contains three input nodes; assume that there are also three output nodes, A, B, C with initial weights randomly chosen:

$$W(0) = \begin{pmatrix} w_1 : & 0.2 & 0.7 & 0.3 \\ w_2 : & 0.1 & 0.1 & 0.9 \\ w_3 : & 1 & 1 & 1 \end{pmatrix} .$$

Assuming $\eta = 0.5$.

Example 5.3 (cont...)

The first sample presented is $i_1 = (1.1, 1.7, 1.8)$. Squared Euclidean distance between A and i_1 :

$$d_{1,1}^2 = (1.1 - 0.2)^2 + (1.7 - 0.7)^2 + (1.8 - 0.3)^2 = 4.1. \text{ Similarly, } d_{2,1}^2 = 4.4 \text{ and } d_{3,1}^2 = 1.1.$$

C is the “winner” since $d_{3,1}^2 < d_{1,1}^2$ and $d_{3,1}^2 < d_{2,1}^2$. A and B are therefore not perturbed by this sample whereas C moves halfway towards the sample (since $\eta = 0.5$).

The resulting weight matrix is:

$$W(1) = \begin{pmatrix} w_1 : & 0.2 & 0.7 & 0.3 \\ w_2 : & 0.1 & 0.1 & 0.9 \\ w_3 : & 1.05 & 1.35 & 1.4 \end{pmatrix}$$

$$\text{E.g. } w_{31} = w_{31} + \Delta w_{31} = 1 + 0.5 * (1.1 - 1.0) = 1.05$$

Example 5.3 (cont...)

Assuming input vectors are presented repeatedly in the sequence, the weight matrix after 12 steps is :

$$W(12) = \begin{pmatrix} w_1 : & 0.55 & 0.3 & 0.3 \\ w_2 : & 0 & 0.4 & 1.35 \\ w_3 : & 1 & 1.2 & 1.25 \end{pmatrix} .$$

- Each node tends to be the winner for the same input vectors, in later iterations, and moves towards their centroid.
- Convergence is not smooth when η is high.
- Different results are observed when other distance measures are used. E.g. use **Manhattan distance** $d(x, y) = \sum |x_i - y_i|$, and you will have different centroids.
- **Using three nodes does not guarantee that three "clusters" will be obtained** by the network. Close initial values may result in splitting of a cluster into two.
- Results depend on the **initial weight vectors and the sequence in which samples are presented** especially when η is high.

k-means Clustering

- A statistical procedure closely related to simple competitive learning, which **computes cluster centroids directly instead of making small updates** to node positions.
- The k-means algorithm is fast and converges to a state in which each prototype changes little, assuming that successive vectors presented to the algorithm are drawn by independent random trials from the input data distribution.
- Both k-means and simple competitive learning can lead to local minima of E .

Algorithm k-means Clustering

Figure 5.5 k-means clustering algorithm

Initialize k prototypes

$$w_j = i_\ell, j \in \{1, \dots, k\}, \ell \in \{1, \dots, P\}$$

Each cluster C_j is associated with prototype w_j .

repeat

for each input vector i_ℓ do

Place i_ℓ in the cluster with nearest prototype w_{j^*}

end for

for each cluster C_j do


$$w_j = \frac{1}{|C_j|} \sum_{i_l \in C_j} i_l \quad \text{where } |C_j| \text{ is the cluster size}$$

end for

Compute $E = \sum_{j=1}^k \sum_{i_l \in C_j} |i_l - w_j|^2$

until E no longer decreases, or cluster memberships stabilize

k-means Clustering



- $k=3$. Initially 3 centroids are initialized by 3 training samples
- Takes only two iterations to stabilize

Kohonen as k-means

The simple competitive learning algorithm conducts stochastic gradient descent on the *quantization error*

$$E = \sum_p |i_p - \mathbb{W}(i_p)|^2$$

where $\mathbb{W}(i_p)$ is the weight vector nearest to i_p

The number of nodes is assumed to be fixed, but the right choice may not be obvious. We may attempt to minimize $E + c(\text{number of nodes})$ instead of E , for $c > 0$.

The above adjustment is called **Regularization** (one approach).

Learning Vector Quantizers (LVQ)

- Sometimes clustering is used as a useful preprocessing step for solving classification problem.
- A LVQ is an application of the above, uses *winner-take-all network* and illustrates how unsupervised learning can be adapted to solve supervised learning.
- **Class membership is known for each training pattern.**
- Learns the codebook vectors.

Learning Vector Quantizers

- Each output node is associated with an arbitrary class label in the beginning.
 - Each node should be finally associated with approximately the number of training data belonging to that class.
- Initial weights are chosen randomly.
- The learning rate decreases with time - helps the network converge to a state in which weight vectors are stable.
 - e.g., $\eta(t) = 1/t$ or $\eta(t) = a[1 - (t/A)]$ where $a > 0$ and $A > 1$ i.e., variable rate \rightarrow decrease more with time.

LVQ (cont...)

- When pattern i from class $C(i)$ is presented to the network, let the winner node $j^* \in C(j^*)$.
- **If this is the correct class** i.e., $C(i)=C(j^*)$, j^* moves closer to i

$$\Delta w_{j^*,l} = \eta(t)(i_{k,l} - w_{j^*,l})$$

- Otherwise j^* moves away from i .

$$\Delta w_{j^*,l} = -\eta(t)(i_{k,l} - w_{j^*,l})$$

LVQ1 Algorithm

Figure 5.6 LVQ1 algorithm

Initialize all weights $\in [0, 1]$

repeat

 Adjust $\eta(t)$;

 for each i_k do

 find node j^* whose weight vector w_{j^*} is closest to i_k ;

 end for

 for $\ell = 1, \dots, n$ do

 if the class label of node j^* equals the desired class of i_k
 then

$$\Delta w_{j^*,l} = \eta(t)(i_{k,l} - w_{j^*,l})$$

 else

$$\Delta w_{j^*,l} = -\eta(t)(i_{k,l} - w_{j^*,l})$$

 end if

 end for

until network converges or computational bounds are exceeded

Example 5.5

$\{i_1 = (1.1, 1.7, 1.8), i_2 = (0, 0, 0), i_3 = (0, 0.5, 1.5), i_4 = (1, 0, 0), i_5 = (0.5, 0.5, 0.5), \text{ and } i_6 = (1, 1, 1)\}$. Assume only the first and last samples come from Class 1, and

$$W(0) = \begin{pmatrix} w_1 : & 0.2 & 0.7 & 0.3 \\ w_2 : & 0.1 & 0.1 & 0.9 \\ w_3 : & 1 & 1 & 1 \end{pmatrix}.$$

w_1 is associated with Class 1, and the other two nodes with Class 2.

Simply because there are twice as many training data associated with class 2 than class 1.

Example 5.5 (cont...)

Let $\eta(t) = 0.5$ until $t = 6$, then $\eta(t) = 0.25$ until $t = 12$ and $\eta(t) = 0.1$ thereafter.

- ① Sample i_1 , winner w_3 (distance 1.07), w_3 changed to (0.95, 0.65, 0.60). \rightarrow But w_3 is class 2 and i_1 is class 1. So $-\Delta w$ applied.

- ② ::::

Associations between input samples and weight vectors stabilize by the second cycle of pattern presentations. Weight vectors continue to change, converging to centroids of associated input vectors in 150 iterations.